

---

# IT496: Introduction to Data Mining

---



## Lecture 16

# Logistic Regression

Arpit Rana  
21<sup>st</sup> September 2023

---

## Classification

---

We want to create programs that make predictions.

- We have been studying *regression*: regressors are programs that predict numeric target values.
- We turn now to *classification*: classifiers predict an object's class from a *finite set of classes*.

### Example

Given a vector of feature values that describe an email, predict whether the email is *spam* or *ham*.

---

## Notation

---

Our notation will be the same that we used for regression:

- $\mathbf{x}$  for an object,  $\mathbf{y}$  for the actual class label,  $\hat{\mathbf{y}}$  for the predicted class label.
- We assume we have a finite set of labels,  $\mathbf{C}$ , one per class.
  - Given an object  $\mathbf{x}$ , our task is to assign one of the labels  $\hat{\mathbf{y}} \in \mathbf{C}$  to the object.
- We will often use integers for the labels.
  - E.g. given an email, a spam filter predicts  $\hat{\mathbf{y}} \in \{0, 1\}$ , where 0 means *ham* and 1 means *spam*.
  - But a classifier should not treat these as continuous, e.g. it should never output 0.5.

---

## Notation

---

Furthermore, where there are more than two labels, we should not assume a relationship between the labels.

- Suppose there are three classes  $\{1, 2, 3\}$ .
- Suppose we are classifying object  $\mathbf{x}$  and we happen to know that its actual class label is  $\mathbf{y=3}$ .
  - One classifier predicts  $\hat{\mathbf{y}}=1$ .
  - Another classifier predicts  $\hat{\mathbf{y}}=2$ .
- Which classifier has done better?

## A Variation of Classification

---

Given an object  $\mathbf{x}$ , a classifier outputs a label,  $\hat{\mathbf{y}} \in \mathbf{C}$ .

- Instead, a classifier could output a probability distribution over the labels  $\mathbf{C}$ .
- For Example,
  - Given an email  $\mathbf{x}$ , a spam filter might output  $\langle 0.2, 0.8 \rangle$  meaning  $P(y = ham | x) = 0.2$  and  $P(y = spam | x) = 0.8$
  - The probabilities must sum to 1.
- We can convert such a classifier into a more traditional one by taking the probability distribution and selecting the class with the highest probability:

$$\arg \max_{\hat{y} \in C} P(\hat{y} | x)$$

## Types of Classification

---

We distinguish three types of classification:

- **Binary classification**, in which there are just two classes, i.e.  $|\mathbf{C}|=2$ , e.g. fail/pass, ham/spam, benign/malignant.
- **Multiclass classification**, where there are more than two classes, i.e.  $|\mathbf{C}|>2$ , e.g. let's say that a post to a forum or discussion board can be *a question, an answer, a clarification or an irrelevance*.
- **Multilabel classification**, where the classifier can assign  $\mathbf{x}$  to more than one class. I.e. it outputs a set of labels,  $\hat{\mathbf{y}} \subseteq \mathbf{C}$ .
  - E.g. consider a movie classifier where the classes are genres, e.g.  $\mathbf{C} = \{\text{comedy, action, horror, documentary, romance, musical}\}$ .
  - The classifier's output for *The Blues Brothers* should be  $\{\text{comedy, action, musical}\}$ .

Do not confuse this with multiclass classification.

---

## Types of Classification

---

In fact, there are even more types of classification, but we will not be studying them further:

- **Ordered classification**, there is an ordering defined on the classes.
  - The ordering matters in measuring the performance of the classifier.
- E.g. consider a classifier that predicts a student's overall grade, i.e. {A, B, C, D}.
  - Suppose for student  $\mathbf{x}$ , the actual class  $\mathbf{y}=\mathbf{A}$ .
  - One classifier predicts  $\hat{\mathbf{y}}=\mathbf{B}$ .
  - Another classifier predicts  $\hat{\mathbf{y}}=\mathbf{C}$ .
  - Which classifier has done better?

## Binary Classification

---

In binary classification, there are two classes.

- It is common to refer to one class (the one labelled 0) as the **negative class** and the other (the one labelled 1) as the **positive class**.
- It doesn't really matter which is which.
  - But, usually, we treat the class we are trying to identify, or the class that requires special action, as the positive class.
  - E.g. in spam filtering, ham is the negative class; spam is the positive class.
  - What about *tumour* classification?
- This terminology is extended to other things too, e.g. we can refer to **negative examples** and **positive examples**.



---

## Class Exercise

---

Consider:

- Predicting tomorrow's rainfall.
- Predicting whether Ireland will have a white Christmas.
- Predicting the sentiment of a tweet (negative, neutral or positive).
- Predicting a person's opinion of a movie on a rating scale of 1 star (rotten) to 5 stars (fab).

Answer the following:

- Which are regression and which classification?
- If classification, which are binary and which are multiclass?
- If binary, which is the positive class and which the negative?

---

# Logistic Regression

## The Model

---

## Logistic Regression

---

We can use *model-based learning* for classification.

- There are all sorts of model, but the simplest again is a linear model.
- For **regression**, we wanted to find the line/plane/hyperplane that best fits the training examples.
- For **classification**, we want to find the line/plane/hyperplane that best separates training examples of different classes.

For two features, if it is possible to find a line that separates the data (only **positive examples** on one side, only **negative examples** on the other), we say the dataset is linearly separable.

This generalizes from straight lines to planes and hyperplanes in the case of more features.

---

## Logistic Regression

---

Despite its name, logistic regression is used for classification.

- At heart, it predicts a number (and turns it into a probability), and perhaps this is why its name mentions regression.
- At heart, it builds linear models.

## Logistic Regression for Binary Classification

Let's start with logistic regression for binary classification.

- In this case, logistic regression predicts the probability that  $\mathbf{x}$  belongs to the positive class.
- This is what logistic regression does:

$$\hat{y} = \begin{cases} 0 & \text{if } P(\hat{y} = 1 | x) < 0.5 \\ 1 & \text{if } P(\hat{y} = 1 | x) \geq 0.5 \end{cases}$$

**Thresholding:** so it outputs one of the two class labels

Where,  $P(\hat{y} = 1 | x) = \sigma(x\beta)$

Where, 
$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

and  $\mathbf{x}\beta$  is familiar from linear regression (and assumes that  $\mathbf{x}$  has an extra 'feature',  $\mathbf{x}_0 = 1$ )

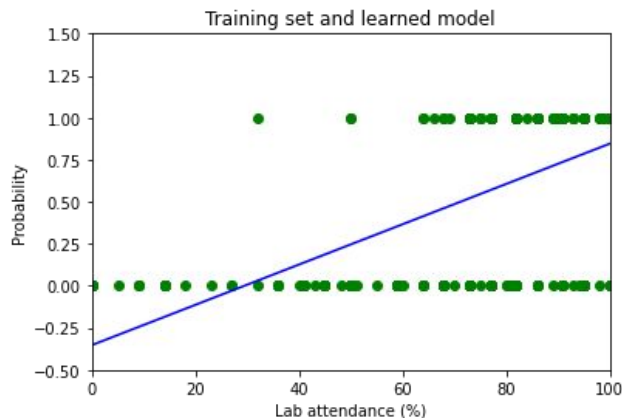
## Why Not Just Linear Regression

In Linear Regression, each hypothesis  $h_{\beta}$  was of the form

$$\begin{aligned} h_{\beta}(x) &= \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_n x_n \\ &= x\beta \end{aligned}$$

where  $\mathbf{x}$  is a row vector with  $(n+1)$  elements (with  $\mathbf{x}_0 = 1$ ), and  $\beta$  is a (column) vector of coefficients.

Why can't we just use this directly to predict probabilities?



Probability of passing  
IT492 course given  
students' lab attendance

# Why Not Just Linear Regression

## The Logistic Function

To 'squash' the values of  $\mathbf{x}\beta$  to  $[0, 1]$  so we can treat them as probabilities

$$h_{\beta}(x) = \sigma(x\beta)$$


where  $\sigma$  is the logistic function  
(also called the 'logit'):

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

The logistic function is also called the **sigmoid function** (which is what we will call it) because it is S-shaped:

A minor point: the sigmoid function asymptotically approaches 0 and 1.




- So, in fact its values are in  $(0, 1)$  and not  $[0, 1]$ .



## Logistic Regression: Loss Function

For a given example  $\mathbf{x}$ , we propose the binary cross-entropy loss as the loss function -

$$J(x, y, \beta) = -(y \log(h_{\beta}(x)) + (1 - y) \log(1 - h_{\beta}(x)))$$





---

## Logistic Regression: Loss Function

---

The overall loss function,  $J$ , is simply the average of this over all the training examples

$$J(X, y, \beta) = -\frac{1}{m} \sum_{i=1}^m \left[ y^{(i)} \log \left( h_{\beta} \left( x^{(i)} \right) \right) + \left( 1 - y^{(i)} \right) \log \left( 1 - h_{\beta} \left( x^{(i)} \right) \right) \right]$$

This loss function is sometimes called the **log loss** function.

## Logistic Regression: Gradient Descent

---

So how do we find the hypothesis that minimizes this log loss function?

- Happily, this function is convex.
- But there is no equivalent to the Normal Equation, so we must use **Gradient Descent**.
- Not that it matters, but here is the partial derivative of its loss function with respect to  $\beta_j$


$$\frac{\partial J(X, y, \beta)}{\partial \beta_j} = \frac{1}{m} \sum_{i=1}^m \left( x^{(i)} \beta_j - y^{(i)} \right) \times x_j^{(i)}$$

Since we'll be using Gradient Descent, we must remember to scale our data.

## Logistic Regression: Decision Boundary


The examples for which logistic regression predicts probabilities of 0.5,  $P(\hat{y} = 1 | x) = 0.5$  lie on what is called the *decision boundary*.

- If you look at the graph of the sigmoid function, its output is 0.5 when its input ( $z$ ) is zero. It follows that the decision boundary are examples where  $\mathbf{x}\beta = 0$ .
- For example, we use the iris dataset that contains the *sepal* and *petal length* and *width* of 150 iris flowers of three different species: *Iris-Setosa*, *Iris-Versicolor*, and *Iris-Virginica*.
- Let's try to build a classifier to detect the *Iris-Virginica* type based only on the *petal width* feature.



## Logistic Regression: Decision Boundary

- Let's try to build a classifier to detect the Iris-Virginica type using the same dataset but this time displaying two features: *petal width* and *length*.



---

## Logistic Regression in Scikit-learn

---

### FYR:

- If you want fine-grained control over the learning rate and so on, then scikit-learn offers you the `SGDClassifier` class.
- But most people use the `LogisticRegression` class, which sits on top of the `SGDClassifier` class.
- If you want to use regularization with Logistic Regression, then there is a separate scikit-learn class for ridge classification (`RidgeClassifier` but *none* for lasso).
- But you can instead use `LogisticRegression`, which has an argument called *penalty*, whose possible values include "l1" and "l2". The amount of regularization is usually controlled by a hyperparameter called *alpha* in the `Lasso` and `Ridge` classes.
- For `LogisticRegression`, this hyperparameter is called **C** and **C** is the inverse of *alpha*, so small values means more regularization!

Next lecture

---

# **Multinomial Logistic Regression**

22<sup>nd</sup> September 2023

---