
IT496: Introduction to Data Mining



Lecture 27-28

Training DNN: Vanishing Gradient Problem

(Slides are created from the lecture notes of Dr. Derek Bridge, UCC, Ireland)

Arpit Rana
20th October 2023

Problems in Training Deep Neural Network

- ***Vanishing Gradient Problem***

The gradients grow smaller and smaller when flowing backward through the DNN while training.

- ***Need enough (labelled) training data***

Large network demands huge amount of data or it might be too costly to label.

- ***Training may be extremely slow***

Training very large DNN using (batch) gradient descent can be painfully slow.

- ***Overfitting***

A model with millions of parameters would severely risk overfitting the training data.

Problems and Solutions in Training Deep Neural Network

- ***Vanishing Gradient Problem***
 - Better Initialization
 - Non-saturating Activation Functions
 - Batch Normalization
- ***Need enough (labelled) training data***
 - Reusing pretrained layers (transfer learning)
 - Unsupervised pre-training
 - Pre-training on an auxiliary task
- ***Training may be extremely slow***
 - Using faster optimizers
 - Learning rate scheduling
- ***Overfitting***
 - Reducing the network size
 - Weight regularization
 - Dropout
 - Early stopping

Problems and Solutions in Training Deep Neural Network

- ***Vanishing Gradient Problem***

- Better Initialization
- Non-saturating Activation Functions
- Batch Normalization

- ***Need enough (labelled) training data***

- Reusing pretrained layers (transfer learning)
- Unsupervised pre-training
- Pre-training on an auxiliary task

- ***Training may be extremely slow***

- Using faster optimizers
- Learning rate scheduling

These will not be covered in this course.

- ***Overfitting***

- Reducing the network size
- Weight regularization
- Dropout
- Early stopping

This will be covered in the next lecture.


Vanishing Gradient Problem

Activation Function: sigmoid (logistic function)

Weight Initialization Technique: random initialization (using a normal distribution with a mean 0 and a standard deviation of 1).

X. Glorot and Y. Bengio found that this combination leads to the following.


- The variance of the outputs of each layer is much greater than the variance of its inputs.
- The variance keeps increasing after each layer until the activation function saturates at the top layers.



Vanishing Gradient Problem

Looking at the logistic function and its derivative plots below.

- When inputs become large (negative or positive), the function saturates at 0 or 1, with a derivative extremely close to 0.



- Thus, during backpropagation, there's not much gradient to propagate back to earlier layers (and what little gradient there is, gets diluted as it propagates back).

Vanishing Gradient Problem

- Each weight is updated by an amount proportional to the gradient of the loss function with respect to that weight.
 - But if the gradient is very small, the weight doesn't change much.
 - This may prevent the network from converging to a good approximation of the target function.
 - We can now see that this is worse for deeper networks.
 - The error signal becomes ever smaller as it is back propagated.
- We look at the following solutions:
 - Better weight initialization;
 - Non-saturating activation functions;
 - Batch normalization.

Weight Initialization

Glorot and Bengio (2010) argued that -

- the variance of the outputs of each layer should be equal to the variance of its inputs (going forward), and
- the gradients should have equal variance before and after flowing through a layer (going backward).

Both the conditions are not guaranteed unless the number of input units to a layer (fan_{in}) equals the number of neurons, i.e., number of output units (fan_{out})

Weight Initialization

- Glorot (Xavier) Initialization: connection weights must be initialized randomly while using the logistic activation function -
 - Normal distribution with mean 0 and variance $\sigma^2 = 1 / fan_{avg}$
 - Or a uniform distribution between $-r$ and $+r$, with $r = \sqrt{3 * fan_{avg}}$
where, $fan_{avg} = (fan_{in} + fan_{out}) / 2$
- Other similar strategies for different activation functions were proposed -

Initialization	Activation functions	σ^2 (Normal)
Glorot	None, Tanh, Logistic, Softmax	$1 / fan_{avg}$
He	ReLU & variants	$2 / fan_{in}$
LeCun	SELU	$1 / fan_{in}$

Weight Initialization in Keras

- By default, Keras uses *Glorot initialization* with a *uniform* distribution.
- You can change this to *He* initialization by setting `kernel_initializer="he_uniform"` or `kernel_initializer="he_normal"` when creating a layer, like this:

```
keras.layers.Dense(10, activation="relu", kernel_initializer="he_normal")
```

Non-Saturating Activation Functions

Certain activation functions, including the sigmoid function, are one cause of the vanishing gradient problem.


- When the input to this function becomes large (positive or negative), the function saturates (i.e. becomes very flat).
 - Even when the gradient is at its greatest (when input z is 0 and $\sigma(z) = 0.5$), it is only 0.25 (gradient = $\sigma(z)(1 - \sigma(z))$).
 - So in the back propagation, gradients always diminish by a quarter or more.
- This is why we rarely use the sigmoid function as the activation function in the hidden layers of deep networks.

Non-Saturating Activation Functions

- Lots of alternatives have been proposed, including the ReLU (*rectified linear unit*) activation function,
- ReLU does not saturate for positive values and it is quite fast to compute.

$$\text{ReLU}(z) = \max(0, z)$$

$$\text{Softplus}(z) = \log(1 + e^z)$$



ReLU and its smooth variant, *softplus*.


Non-Saturating Activation Functions

- Neurons that use the ReLU activation function have obvious problems too:
 - If their input (weighted sum) is negative, the output is zero; and the gradient is zero; and
 - if this is true for all examples in the training set then, in effect, the neuron dies. This is known as “*dying* ReLU” problem.
 - The gradient changes abruptly at $z = 0$, which can make Gradient Descent bounce around.
- Despite its problems, it remains a popular choice.

Non-Saturating Activation Functions

Leaky ReLU have been proposed having at least some non-zero gradient for negative inputs (alternatives to ReLU are slower to compute and they introduce further hyperparameters.)

The hyperparameter α defines how much the function “leaks”: it is the slope of the function for $z < 0$, and is typically set to 0.01.



```
layers.Dense(7 * 7 * 128),  
layers.LeakyReLU(alpha=0.2),
```

Instead of being a hyperparameter, α becomes a parameter that can be modified by backpropagation like any other parameter, this version of ReLU is known as *Parametric Leaky ReLU (PReLU)*.


Non-Saturating Activation Functions

ELU (Exponential Linear Unit)

- It helps alleviate the vanishing gradient problem.
- It has a non-zero gradient for $z < 0$, which avoids the dead neurons problem.
- If $\alpha = 1$, then the function is smooth everywhere, which helps speed up gradient descent (does not bounce much left and right of $z = 0$.)

$$ELU_{\alpha}(z) = \begin{cases} \alpha(\exp(z) - 1) & \text{if } z < 0 \\ z & \text{if } z \geq 0 \end{cases}$$

The hyperparameter α defines the value that the ELU function approaches when z is a large negative number.



Non-Saturating Activation Functions

SELU (Scaled Exponential Linear Unit)

- If we build a neural network with a stack of dense layers, and all the hidden layers use SELU activation, then the network will “*self-normalize*” (the output of each layer will tend to preserve mean 0 and standard deviation 1 during training).
- The input features must be standardized (mean 0 and standard deviation 1).
- Every hidden layer weights must also be initialized using LeCun normal initialization. In Keras, this means setting `kernel_initializer="lecun_normal"`.

$$SELU(z) = \begin{cases} \lambda \alpha (\exp(z) - 1) & \text{if } z < 0 \\ \lambda z & \text{if } z \geq 0 \end{cases} \quad [\alpha \text{ (>=1) and } \lambda \text{ (>1) are predefined constants.}]$$

[illegible]

Non-Saturating Activation Functions

Which activation function should we use for the hidden layers of our deep neural network?

- In general, **SELU** > **ELU** > *leaky ReLU (and its variants)* > **ReLU** > *tanh* > *logistic*
- If the network's architecture prevents it from self-normalizing, then ELU may perform better than SELU (since SELU is not smooth at $z = 0$).
- If you care a lot about runtime latency, then you may prefer leaky ReLU. If you don't want to tweak yet another hyperparameter, you may just use the default α values used by Keras (e.g., 0.3 for the leaky ReLU).
- If you have spare time and computing power, you can use cross-validation to evaluate other activation functions, in particular PReLU if you have a huge training set.

Batch Normalization

We previously studied the usefulness of feature scaling when doing Gradient Descent for, e.g., linear regression.

... and we've been doing this to the features in our neural networks too.

- But, if this is a good idea for the inputs to the first hidden layer, why not use the same idea for the inputs to subsequent layers?
- In other words, we normalize the activations (outputs) of layer l prior to them being used as inputs to layer $l + 1$.
- This will control the distribution of the values throughout the training process.
- This, in essence, is the idea of *batch normalization*.

Batch Normalization

In summary, for a given layer, it standardizes the outputs of the neurons (subtract the mean, divide by the standard deviation), *then it scales the result and adds an offset.*

Batch Normalization Algorithm:

$$1. \quad \mu_B = \frac{1}{m_B} \sum_{i=1}^{m_B} \mathbf{x}^{(i)}$$

$$2. \quad \sigma_B^2 = \frac{1}{m_B} \sum_{i=1}^{m_B} \left(\mathbf{x}^{(i)} - \mu_B \right)^2$$

$$3. \quad \hat{\mathbf{x}}^{(i)} = \frac{\mathbf{x}^{(i)} - \mu_B}{\sqrt{\sigma_B^2 + \epsilon}}$$

$$4. \quad \mathbf{z}^{(i)} = \gamma \otimes \hat{\mathbf{x}}^{(i)} + \beta$$

During training, four parameter vectors are learned in each batch-normalized layer:

- γ (the output scale vector) and β (the output offset vector) are learned through regular backpropagation, and
- μ (the final input mean vector), and σ (the final input standard deviation vector) are estimated using an exponential moving average.

What about the test time?

Batch Normalization

- Batch Normalization reduces the vanishing gradients problem so much, we can even use saturating activation functions.
- Training becomes less sensitive to the method used for randomly initializing weights.
- Much larger learning rates work (faster convergence) with less risk of divergence.
- It acts like a regularizer.

Batch Normalization in Keras

- Just add another layer!
- For example,
 - `x = Dense(512, activation="relu")(x)`
 - `x = BatchNormalization()(x)`
- This is how we will do batch normalization in this course.
- Ignore: in fact, there is some debate about whether we should batch normalize the activations of the layer (as above) or the weighted sum, before applying the activation function (as below):
 - `x = Dense(512, activation="linear", use_bias=False)(x)`
 - `x = BatchNormalization()(x)`
 - `x = Activation("relu")(x)`
- We will stick with the former, which is more concise.

Batch Normalization in Keras

```
model = keras.models.Sequential([
    keras.layers.Flatten(input_shape=[28, 28]),
    keras.layers.BatchNormalization(),
    keras.layers.Dense(300, activation="elu", kernel_initializer="he_normal"),
    keras.layers.BatchNormalization(),
    keras.layers.Dense(100, activation="elu", kernel_initializer="he_normal"),
    keras.layers.BatchNormalization(),
    keras.layers.Dense(10, activation="softmax")
])
```

Layer (type)	Output Shape	Param #
=====		
flatten_3 (Flatten)	(None, 784)	0
=====		
batch_normalization_v2 (Batch Normalization)	(None, 784)	3136
=====		
dense_50 (Dense)	(None, 300)	235500
=====		
batch_normalization_v2_1 (Batch Normalization)	(None, 300)	1200
=====		
dense_51 (Dense)	(None, 100)	30100
=====		
batch_normalization_v2_2 (Batch Normalization)	(None, 100)	400
=====		
dense_52 (Dense)	(None, 10)	1010
=====		
Total params: 271,346		
Trainable params: 268,978		
Non-trainable params: 2,368		

Exploring Vanishing Gradient with MNIST

```
def build_mnist_network(activation, initializer, use_batch_norm):
    inputs = Input(shape=(28 * 28,))
    x = Rescaling(scale=1./255)(inputs)
    x = Dense(512, activation=activation, kernel_initializer=initializer)(x)
    if use_batch_norm:
        x = BatchNormalization()(x)
    outputs = Dense(10, activation="softmax", kernel_initializer=initializer)(x)
    model = Model(inputs, outputs)
    model.compile(optimizer=RMSprop(learning_rate=0.0001), loss="sparse_categorical_crossentropy", metrics=["accuracy"])
    return model
```

```
networks = [
    build_mnist_network("sigmoid", "random_normal", False),
    build_mnist_network("sigmoid", "random_normal", True),
    build_mnist_network("sigmoid", "glorot_uniform", False),
    build_mnist_network("sigmoid", "glorot_uniform", True),
    build_mnist_network("relu", "random_normal", False),
    build_mnist_network("relu", "random_normal", True),
    build_mnist_network("relu", "glorot_uniform", False),
    build_mnist_network("relu", "glorot_uniform", True)
]

for network in networks:
    network.fit(mnist_x_train, mnist_y_train, epochs=10, batch_size=32, verbose=0)
    test_loss, test_acc = network.evaluate(mnist_x_test, mnist_y_test, verbose=0)
    print(test_acc)
```

Output:

0.9451000094413757
0.9549999833106995
0.9435999989509583
0.9496999979019165
0.9746999740600586
0.9814000129699707
0.9746999740600586
0.9840999841690063

Hyperparameter Tuning

- First an observation about validation sets:
 - When using neural networks, we typically have a large dataset.
 - Hence, we use holdout to split the dataset into train, validation and test sets.
 - If you have a smaller dataset, where you can afford to split off a test set, but you cannot afford to split off a validation set, then you might want to use k-fold cross-validation, as we did in our scikit-learn examples.
 - Keras does not have any in-built cross-validation functions. You'd have to write your own.

Hyperparameter Tuning

- Second, an observation about hyperparameter tuning:
 - Hyperparameter tuning involves training lots of different models with different values for the hyperparameters, and comparing their performance on the validation data.
 - This is often a problem with neural networks: training can be slow (due to large datasets and models that have many parameters) and so training lots of different models for hyperparameter tuning is very time-consuming.
 - People often don't bother! They just pick hyperparameter values out of their heads, or accept the Keras defaults.
 - If you can afford to be more systematic, there is a separate library (which you would need to install) called KerasTuner. It automates hyperparameter tuning, similar to what we saw in scikit-learn.

Next lecture

Overfitting in Neural Network

20th October 2023
